

Vague memories need a recovery path—not another search box.

MemoryTrail helps people keep clues, question guesses and explore nearby moments.

THE USER'S JOB

**“Find the café with blue chairs.
I think it was in Goa.”**

CERTAIN · blue chairs

MAYBE · Goa

Illustrative scenario—not an interview quote.

The opportunity to test

Make uncertain clues visible and reversible, then let a near-match become a path to the target.

The baseline is already conversational.

Ask Photos has natural language and follow-ups.
Test recovery—not chatbot novelty.

12 public excerpt records

0 user interviews completed

Public deployment pending

Successful retrieval is a funnel, not a search-count metric.

North star: verified retrieval users / eligible users attempting a partial-memory retrieval.

N · START

Eligible attempt

A remembered target is known to be in the collection.

C · CLUES

Usable clues

The user supplies or selects usable memory clues.

V · CANDIDATE

Relevant photo seen

Ground truth confirms a relevant candidate was displayed.

O · OPEN

Target opened

The intended photo is opened—not merely any result.

S · VERIFIED

Correct retrieval

The intended target is confirmed within the task window.

$$S / N = (C / N) \times (V / C) \times (O / V) \times (S / O)$$

Count each eligible user once on their first standardized task in a 7-day window. Proposed task limit: 120 seconds. Events are nested; zero-denominator rates are undefined.

A repeatable workflow separates reported experiences from hypotheses.

Evidence Lab checks source-linked reports and keeps human review visible.

01

Import & deduplicate

Source URL + exact text.
Up to 500 records.

02

Extract with AI

AI returns a stage, exact quote and interpretation.

03

Validate evidence

Check IDs and exact quotes.
Reject invalid output.

04

Review the source

Check context, relevance and counterevidence.

05

Compare & export

Compare stage groups and export linked evidence.

Seed: 12 excerpts / 5 threads / 3 communities

A small desk-research seed—not at-scale validation.

Live extraction: batches of 20.
Provider calls and public URL: pending.

Public reports expose several failures—not one validated root cause.

The seed includes historical complaints, positive experiences and one adjacent scoping issue.

RETRIEVE • R02

Combined clues can miss

One poster reports no results for a season + city + year query. This may be an indexing or matching issue.

REFINE • R03

A near-match can become a dead end

A poster reports losing navigation from a result to other photos from the same date.

RETRIEVE • R05 / R09

Candidate coverage can be incomplete

Bird and cat searches are reported to return only a small subset. Better recovery cannot fix a missing index.

COUNTEREVIDENCE • R04 / R10

Search also works for some people

Other posters report useful results or improved retrieval. Do not infer a universal product failure.

Start with travellers who remember a scene, not exact dates.

Provisional problem: uncertain clues may block recovery of a photo that is actually present.

PROVISIONAL SEGMENT

Scene remembered; metadata uncertain

Recruit Google Photos users with 2+ years of trip history and a recent difficult retrieval attempt.

ROOT-CAUSE HYPOTHESIS

Guesses may derail the recovery path

Users may need to separate known clues from guesses, then explore the event around a plausible result.

REQUIRED FIELD RESEARCH

Observe 5–6 users before committing

Reconstruct a recent incident. Watch an unprimed retrieval task. Record where the search actually breaks.

Pivot if the target is not indexed—or current Ask Photos already solves the observed task.

Expose uncertainty before asking users to rewrite the search.

Three reversible steps test the recovery hypothesis without treating a guess as a hard filter.

01 · KEEP THE MEMORY


Separate known from “maybe”

CERTAIN · blue chairs

MAYBE · Goa

Edit, ignore or restore any clue.
Guesses do not exclude images.

02 · INSPECT CANDIDATES



Matches: blue chairs

Illustration · synthetic metadata

03 · RECOVER THE MOMENT

Explore nearby, then confirm

Open photos from the same known event. Reject a wrong match, undo, or answer one clarifying question.

NOT SURE is always valid

The MVP preserves alternatives and makes recovery reversible.

A self-contained demo runs now; the server-side AI path still requires live verification.

RUNS WITHOUT SETUP

Illustrated retrieval + Evidence Lab

Clue editing · soft guesses · reject / undo ·
known-event traversal · source review · session
export

IMPLEMENTED, NOT LIVE-TESTED

Selected images → Gemini → candidates

Image captioning, structured clues and
reranking. Private images are selected explicitly;
no full-library integration.

22

unit tests

including mocked provider
output

11

browser checks

standalone interaction flows

7

HTTP smoke checks

local server; no model key

Three user tests must challenge the hypothesis, not confirm it.

Return to at least 3 target-segment users. Completed participant tests: 0.

1 • Observe the baseline

Watch a real Google Photos attempt first. Record the account / version, target availability and failure point.

2 • Test the mechanism

Use a consented, controlled photo subset. Rotate different matched targets and counterbalance conditions.

3 • Verify the outcome

Verify the target independently. Record time, errors, recovery actions and the user's explanation.

PROPOSED PILOT GATE

**2 of 3 unaided
0 false verified
successes**

A usability checkpoint, not statistical evidence of uplift.

The offline keyword toggle is a toy baseline—not Google Photos. Do not report Google Photos uplift from that comparison.

Measure verified recovery while guarding trust and privacy.

No baseline or participant outcome has been measured. Thresholds below are proposed.

METRIC LAYER	DEFINITION / USE	CURRENT STATUS
User outcome	Verified users retrieving the correct target within 120 seconds / eligible users (first task).	No baseline. Proposed task limit: 120 seconds.
Task diagnostics	Task success; time across all attempts; abandonment; clue edits; near-match recovery.	Record failed and uncertain attempts—not just clicks.
AI / system quality	Candidate coverage; wrong interpretations; latency; calls and cost per attempt.	Live quality and cost checks are pending.
Trust guardrails	False verified success; image exposure; misleading metadata; inability to undo.	Proposed goal: no false confirmations or exposure.

A click is not proof. Real-photo selections remain unverified until the intended target is independently checked.

Publish only after evidence, privacy and access checks pass.

Fieldwork, live model checks and public testable links are required before submission.

RISKS → MITIGATIONS

Wrong or absent memories

Keep guesses soft; check collection coverage; never manufacture a missing photograph.

Private photos or uncontrolled cost

Explicit selection and consent; server-side keys; reviewer access; provider spending limits.

No advantage over current search

Compare against observed Ask Photos behaviour; pivot when evidence rejects the recovery hypothesis.

SUBMISSION CHECK

BUILT

Deck + local demo
AI code + research templates

PENDING

5–6 interviews; 3 user tests
Live AI and public URLs
Replace draft with findings

Public URLs are intentionally not invented.

Brief deadline shown: 7 October 2026, 4:00 PM Asia/Calcutta. Verify the portal before submitting.